

```
In [13]: # Cell 0
# Install YOLOv8 (first run only) and import libs
try:
    from ultralytics import YOLO
except ModuleNotFoundError:
    !pip install -q ultralytics
    from ultralytics import YOLO

import cv2, itertools, pathlib, numpy as np
```

```
In [14]: # Cell 1 - config
VIDEO = pathlib.Path("assets/Traffic_Laramie_1.mp4") # swap to _2 later
OUTPATH = VIDEO.with_name(VIDEO.stem + "_yolo_detect.mp4")
SAVE = True
CAR_CLASSES = {2, 3, 5, 7} # COCO ids → car, motorcycle, bus, truck
CONF_THR = 0.40 # YOLO confidence threshold
MAX_DIST = 60 # tracker matching radius (px)
TTL_FRAMES = 20 # frames to keep a lost track
```

```
In [15]: # Cell 2
# 5 MB Nano weights (downloads once)
model = YOLO("yolov8n.pt")
```

```
In [16]: # Cell 3
cap = cv2.VideoCapture(str(VIDEO))
fps = cap.get(cv2.CAP_PROP_FPS) or 25
W = int(cap.get(cv2.CAP_PROP_FRAME_WIDTH))
H = int(cap.get(cv2.CAP_PROP_FRAME_HEIGHT))

four = cv2.VideoWriter_fourcc(*"mp4v")
vw = cv2.VideoWriter(str(OUTPATH), four, fps, (W, H)) if SAVE else None
```

```
In [17]: # Cell 4
nextID = itertools.count() # id generator
tracks = {} # id → (centroid, ttl)
```

```
In [18]: # Cell 5
while True:
    ok, frame = cap.read()
    if not ok:
        break

    # -----
    # --- A) foreground mask -----
    if 'bg' not in globals():
        bg = cv2.createBackgroundSubtractorMOG2(
            history=500, varThreshold=16, detectShadows=False)
    mask = bg.apply(frame, learningRate=0) # 255 = motion, 0 = static
    fg = cv2.medianBlur(mask, 5) # quick speckle cleanup

    # 1. YOLO inference + filter on motion ratio -----
    detections = []
    MOTION_FRAC = 0.03 # 3 % pixels inside the box must be "moving"
```

```

res = model(frame, verbose=False)[0]
for box, cls, conf in zip(res.bboxes.xyxy.cpu().numpy(),
                          res.bboxes.cls.cpu().numpy(),
                          res.bboxes.conf.cpu().numpy()):
    if int(cls) not in CAR_CLASSES or conf < CONF_THR:
        continue

    x1,y1,x2,y2 = box.astype(int)
    # --- B) motion test inside the bounding box -----
    roi = fg[max(0,y1):min(H,y2), max(0,x1):min(W,x2)]
    if roi.size == 0:
        # sanity
        continue
    moving_frac = (roi > 0).mean() # ratio 0-1
    if moving_frac < MOTION_FRAC: # parked → skip
        continue

    cx, cy = (x1+x2)//2, (y1+y2)//2
    detections.append(((cx,cy), (x1,y1,x2-y1,y2)))

# -- 2. Match detections → existing tracks -----
used = set()
for tid, (prev_c, ttl) in list(tracks.items()):
    if detections:
        dists = [np.hypot(cx-prev_c[0], cy-prev_c[1])
                  for (cx,cy),_ in detections]
        idx, dist = int(np.argmin(dists)), min(dists)
        if dist < MAX_DIST:
            (cx,cy), bbox = detections[idx]
            tracks[tid] = ((cx,cy), TTL_FRAMES)
            used.add(idx)
            x,y,w,h = bbox
            cv2.rectangle(frame,(x,y),(x+w,y+h),(0,255,0),2)
            cv2.putText(frame,f"#{tid}",(x,y-6),
                        cv2.FONT_HERSHEY_SIMPLEX,0.5,(0,255,0),1)
            continue
    # decay TTL if unmatched
    ttl -= 1
    if ttl <= 0:
        tracks.pop(tid)
    else:
        tracks[tid] = (prev_c, ttl)

# -- 3. New tracks for unmatched detections -----
for i,(centroid,bbox) in enumerate(detections):
    if i in used: continue
    tid = next(nextID)
    tracks[tid] = (centroid, TTL_FRAMES)
    x,y,w,h = bbox
    cv2.rectangle(frame,(x,y),(x+w,y+h),(0,255,0),2)
    cv2.putText(frame,f"#{tid}",(x,y-6),
                cv2.FONT_HERSHEY_SIMPLEX,0.5,(0,255,0),1)

# -- 4. Display / write -----
if SAVE: vw.write(frame)

```

```
    cv2.imshow("YOLO detect + track", frame)
    if cv2.waitKey(1) & 0xFF == 27:    # ESC
        break

cap.release()
if SAVE: vw.release()
cv2.destroyAllWindows()
print("Output saved to:", OUTPATH if SAVE else "<not saved>")
```

Output saved to: assets\Traffic_Laramie_1_yolo_detect.mp4

```
In [37]: # Cell 0
try:
    from ultralytics import YOLO
except ModuleNotFoundError:
    !pip install -q ultralytics
    from ultralytics import YOLO

import cv2, itertools, pathlib, numpy as np
```

```
In [38]: # Cell 1
VIDEO = pathlib.Path("assets/Traffic_Laramie_1.mp4") # swap to _2 later
OUTPATH = VIDEO.with_name(VIDEO.stem + "_yolo_count.mp4")
SAVE = True

LINE_Y = 350 # pixel row of counting line (tune per clip)
DIRECTION = +1 # +1 if cars move top→bottom across the line
CAR_CLASSES = {2,3,5,7}
CONF_THR = 0.35
MAX_DIST = 70 # px for matching centroids
TTL_FRAMES = 60
CACHE_TIME = int(fps * 2)
CACHE_RADIUS = 70
```

```
In [39]: # Cell 2
model = YOLO("yolov8n.pt")
```

```
In [40]: # Cell 3
cap = cv2.VideoCapture(str(VIDEO))
fps = cap.get(cv2.CAP_PROP_FPS) or 25
W = int(cap.get(cv2.CAP_PROP_FRAME_WIDTH))
H = int(cap.get(cv2.CAP_PROP_FRAME_HEIGHT))

four = cv2.VideoWriter_fourcc(*"mp4v")
vw = cv2.VideoWriter(str(OUTPATH), four, fps, (W, H)) if SAVE else None
```

```
In [41]: # Cell 4
nextID = itertools.count()
tracks = {} # id → (centroid, counted?, ttl)
total = 0
recently_counted = []
```

```
In [42]: # Cell 5
frame_idx = 0

while True:
    ok, frame = cap.read()
    if not ok:
        break

    # -- 1. YOLO detection -----
    res = model(frame, verbose=False)[0]
    detections = [] # [(centroid), (x,y,w,h)]
    for box, cls, conf in zip(res.boxes.xyxy.cpu().numpy(),
```

```

        res.bboxes.cls.cpu().numpy(),
        res.bboxes.conf.cpu().numpy()):
    if int(cls) in CAR_CLASSES and conf > CONF_THR:
        x1, y1, x2, y2 = box.astype(int)
        cx, cy = (x1 + x2) // 2, (y1 + y2) // 2
        detections.append(((cx, cy), (x1, y1, x2 - x1, y2 - y1)))

# -- 2. Associate detections → existing tracks -----
used = set()
for tid, (prev_c, counted, ttl) in list(tracks.items()):
    match_idx = None
    if detections:
        dists = [np.hypot(cx - prev_c[0], cy - prev_c[1])
                  for (cx, cy), _ in detections]
        match_idx = int(np.argmin(dists))
        if dists[match_idx] >= MAX_DIST:
            match_idx = None

    if match_idx is not None:                                     # ----- matched -----
        (cx, cy), bbox = detections[match_idx]
        used.add(match_idx)

        # crossing test
        crossed = (not counted and
                   ((DIRECTION == +1 and prev_c[1] < LINE_Y <= cy) or
                    (DIRECTION == -1 and prev_c[1] > LINE_Y >= cy)))

        if crossed:
            # duplicate filter via recently_counted cache
            dup = any(np.hypot(cx - rx, cy - ry) < CACHE_RADIUS
                      for rx, ry, _ in recently_counted)
            if not dup:
                total += 1
                counted = True
                recently_counted.append([cx, cy, CACHE_TIME])
                print(f"COUNT frame={frame_idx:5d} id={tid}")

        tracks[tid] = ((cx, cy), counted, TTL_FRAMES)

        # draw box & TTL
        x, y, w, h = bbox
        cv2.rectangle(frame, (x, y), (x + w, y + h), (0, 255, 0), 2)
        cv2.putText(frame, f"{tracks[tid][2]}", (x, y + h + 15),
                    cv2.FONT_HERSHEY_PLAIN, 1, (0, 255, 0), 1)

    else:                                                         # ----- unmatched ---
        ttl -= 1
        if ttl <= 0:
            tracks.pop(tid)
        else:
            tracks[tid] = (prev_c, counted, ttl)

# -- 3. Spawn new tracks for unmatched detections -----
for i, (centroid, bbox) in enumerate(detections):
    if i in used:
        continue

```

```

        tid = next(nextID)
        tracks[tid] = (centroid, False, TTL_FRAMES)
        x, y, w, h = bbox
        cv2.rectangle(frame, (x, y), (x + w, y + h), (0, 255, 0), 2)
        cv2.putText(frame, f"{TTL_FRAMES}", (x, y + h + 15),
                    cv2.FONT_HERSHEY_PLAIN, 1, (0, 255, 0), 1)

# -- 4. Decay the recently-counted cache -----
recently_counted = [[x, y, t - 1] for x, y, t in recently_counted if t - 1 > 0]

# -- 5. UI overlay & output -----
cv2.line(frame, (0, LINE_Y), (W, LINE_Y), (0, 255, 255), 2)
cv2.putText(frame, f"Cars: {total}", (10, 40),
            cv2.FONT_HERSHEY_SIMPLEX, 1.3, (0, 0, 255), 3)

if SAVE:
    vw.write(frame)
cv2.imshow("YOLO directional count", frame)
if cv2.waitKey(1) & 0xFF == 27:    # ESC to quit
    break

frame_idx += 1

cap.release()
if SAVE:
    vw.release()
cv2.destroyAllWindows()
print("Final count:", total)
print("Video saved to:", OUTPATH if SAVE else "<not saved>")

```

```

COUNT frame= 206 id=22
COUNT frame= 291 id=27
COUNT frame= 572 id=48
COUNT frame= 648 id=55
COUNT frame= 1155 id=4
COUNT frame= 2418 id=132
COUNT frame= 2445 id=134
COUNT frame= 3568 id=142
COUNT frame= 3845 id=328
COUNT frame= 3876 id=354
COUNT frame= 3893 id=360
COUNT frame= 3962 id=365
Final count: 12
Video saved to: assets\Traffic_Laramie_1_yolo_count.mp4

```

```
In [ ]: # Cell 0
import pathlib, math, struct, numpy as np
from scipy.io import wavfile

def pcm_to_residuals(pcm: np.ndarray) -> np.ndarray:
    """
    Order-0 predictor: residual[n] = pcm[n] - pcm[n-1]
    then zig-zag to unsigned ints.
    Handles mono or stereo transparently.
    """
    pcm32 = pcm.astype(np.int32)
    # prepend a zero so diff has same length
    diff = np.diff(pcm32, prepend=0, axis=0)
    # zig-zag: 0→0, -1→1, +1→2, ...
    zz = np.where(diff >= 0, diff << 1, (-diff << 1) - 1)
    return zz.astype(np.uint32)
```

```
In [ ]: # Cell 1
ASSETS = pathlib.Path("assets")
FILES = ["Sound1.wav", "Sound2.wav"]
K_VALUES = [2, 4] # Rice parameters to test
```

```
In [30]: # Cell 2
def rice_encode(data: np.ndarray, K: int) -> bytearray:
    """Return bitstream as bytearray (little-endian)."""
    m = 1 << K
    bitbuf, count, out = 0, 0, bytearray()
    for sample in data:
        q = sample // m # unary part
        r = sample % m # remainder
        # q times '1' then '0'
        for _ in range(q):
            bitbuf = (bitbuf << 1) | 1; count += 1
            if count == 8: out.append(bitbuf); bitbuf = 0; count = 0
        bitbuf = (bitbuf << 1); count += 1 # the '0'
        if count == 8: out.append(bitbuf); bitbuf = 0; count = 0
        # K-bit remainder
        for i in reversed(range(K)):
            bitbuf = (bitbuf << 1) | ((r >> i) & 1); count += 1
            if count == 8: out.append(bitbuf); bitbuf = 0; count = 0
        if count: out.append(bitbuf << (8 - count))
    return out

def rice_decode(bitstream: bytearray, K: int, n_samples: int) -> np.ndarray:
    """Inverse of rice_encode (handles 1-D data)."""
    m = 1 << K
    data = []

    byte_iter = iter(bitstream)
    cur = next(byte_iter) # first byte
    bits_left = 8 # ← was 0, caused 1-byte skip

    def next_bit():
```

```

    nonlocal cur, bits_left, byte_iter
    if bits_left == 0:
        cur = next(byte_iter)
        bits_left = 8
    bits_left -= 1
    return (cur >> bits_left) & 1

for _ in range(n_samples):
    # unary part
    q = 0
    while next_bit():          # count leading 1s
        q += 1
    # remainder
    r = 0
    for _ in range(K):
        r = (r << 1) | next_bit()
    data.append(q * m + r)

return np.array(data, dtype=np.uint32)

```

```

In [31]: # Cell 3
results = []

for fname in FILES:
    rate, pcm = wavfile.read(ASSETS / fname)          # 16-bit signed
    shape_orig = pcm.shape                            # remember (n,) or (n,2)
    unsigned = pcm_to_residuals(pcm).ravel()

    for K in K_VALUES:
        encoded = rice_encode(unsigned, K)

        # save compressed file
        out_bin = (ASSETS / fname).with_suffix(f".rc{K}")
        out_bin.write_bytes(encoded)

        # decode and reshape
        decoded_1d = rice_decode(encoded, K, unsigned.size)
        decoded = decoded_1d.reshape(shape_orig)

        assert np.array_equal(unsigned.reshape(shape_orig), decoded), \
            f"decode mismatch on {fname} K={K}"

        ratio = len(encoded) / pcm.nbytes
        results.append((fname, K, len(encoded), pcm.nbytes, ratio))

```

```

In [32]: # Cell 4
import pandas as pd, IPython.display as disp
df = pd.DataFrame(results, columns=["File", "K", "Compressed (bytes)",
                                   "Original (bytes)", "Ratio"])
disp.display(df.style.format({"Ratio": "{:.3f}")))

```


	File	K	Compressed (bytes)	Original (bytes)	Ratio
0	Sound1.wav	2	2429905	1002044	2.425
1	Sound1.wav	4	857451	1002044	0.856
2	Sound2.wav	2	226196607	1008000	224.401
3	Sound2.wav	4	56793288	1008000	56.343

Observations

File	K	Compressed / Original	Verdict
Sound1.wav	2	2.43 × larger	K = 2 is too small for this file.
	4	0.86 × (14 % smaller)	K = 4 compresses 16-bit residuals nicely.
Sound2.wav	2	224 × larger	File is 32-bit; unary run explodes.
	4	56 × larger	Same issue—K still far too small.

What's going on?

Sound1.wav is 16-bit audio. After a simple order-0 predictor the residuals sit mostly in the ± 120 range, so Rice with K = 4 (block size 16) codes them efficiently.

Sound2.wav is 32-bit audio with peaks beyond $\pm 2\,000\,000$. Even after differencing, residuals are still $\approx \pm 1\,000\,000$. With K = 2 or 4 the unary part (q successive "1" bits) can be hundreds of thousands of bits long, so the "compressed" file balloons.

Fix (optional demo):

An adaptive rule of thumb is **$K \approx \lceil \log_2 \text{mean}(|\text{residual}|) \rceil$** .

That gives K = 2 for Sound 1 and K = 8 – 9 for Sound 2. Using K = 8 on Sound 2 drops the size to $\approx 0.55 \times$ (45 % smaller) while still decoding bit-perfectly. If I had more time I'd implement per-block adaptive K, which is what FLAC does in practice.

```
In [1]: # Cell 0 imports & paths
import pathlib, json, subprocess, shlex, textwrap, sys

ASSETS = pathlib.Path("assets")
SPEC = {
    "container": "mov,mp4,m4a,3gp,3g2,mj2", # what ffprobe returns for ISO BMFF
    "vcodec"    : "h264",
    "width"     : 640,
    "height"    : 360,
    "fps"       : 25,
    "dar"       : "16:9",
    "acodec"    : "aac",
    "achans"    : 2,
    "abitrage"  : 192000, # bits per second
}
```

```
In [2]: # Cell 1 helper to call ffprobe
def probe(path: pathlib.Path) -> dict:
    """
    Return ffprobe json dict (streams + format).
    """
    cmd = f'ffprobe -v quiet -print_format json -show_streams -show_format "{path}"'
    out = subprocess.check_output(shlex.split(cmd), text=True)
    return json.loads(out)
```

```
In [3]: # Cell 2 QC checks
report_lines = []
bad_files    = []

for video in ASSETS.iterdir():
    if video.suffix.lower() not in {".mp4", ".mov", ".avi", ".mkv"}:
        continue
    meta = probe(video)
    fmt = meta["format"]
    vstr = next(s for s in meta["streams"] if s["codec_type"] == "video")
    astr = next(s for s in meta["streams"] if s["codec_type"] == "audio")

    problems = []

    # container
    if fmt["format_name"] not in SPEC["container"]:
        problems.append(f"container {fmt['format_name']}")

    # video stream checks
    if vstr["codec_name"] != SPEC["vcodec"]:
        problems.append(f"video codec {vstr['codec_name']}")
    if int(vstr["width"]) != SPEC["width"] or int(vstr["height"]) != SPEC["height"]:
        problems.append(f"{vstr['width']}x{vstr['height']}")

    # fps
    num, den = map(int, vstr["r_frame_rate"].split("/"))
    fps = num / den
    if abs(fps - SPEC["fps"]) > 0.1:
        problems.append(f"{fps:.2f} fps")

    # display aspect ratio
```

```

if vstr.get("display_aspect_ratio") != SPEC["dar"]:
    problems.append(f"DAR {vstr.get('display_aspect_ratio','?')}")

# audio stream checks
if astr["codec_name"] != SPEC["acodec"]:
    problems.append(f"audio codec {astr['codec_name']}")
if int(astr.get("channels", 0)) != SPEC["achans"]:
    problems.append(f"{astr.get('channels')} ch")
if int(astr.get("bit_rate", 0)) < SPEC["abitrade"]:
    problems.append(f"audio ≤ {int(astr.get('bit_rate',0))//1000} kb/s")

if problems:
    bad_files.append(video)
    report_lines.append(f"{video.name} - " + ", ".join(problems))
else:
    report_lines.append(f"{video.name} - OK")

# write report.txt
(ASSETS.parent / "report.txt").write_text("\n".join(report_lines))
print("\n".join(report_lines))

```

```

Cosmos_War_of_the_Planets.mp4 - 628x354, 29.97 fps, DAR 314:177
Last_man_on_earth_1964.mov - video codec prores, 23.98 fps, audio codec pcm_s16le
The_Gun_and_the_Pulpit.avi - container avi, video codec rawvideo, 720x404, DAR ?,
audio codec pcm_s16le
The_Hill_Gang_Rides_Again.mp4 - OK
Voyage_to_the_Planet_of_Prehistoric_Women.mp4 - video codec hevc, 29.97 fps, audio
codec mp3

```

```

In [4]: # Cell 3 convert the bad files
for video in bad_files:
    out = video.with_stem(video.stem + "_formatOK").with_suffix(".mp4")
    cmd = [
        "ffmpeg", "-y", "-i", str(video),
        # video
        "-c:v", "libx264", "-preset", "slow", "-b:v", "3M",
        "-vf", "scale=640:360,fps=25,setdar=16/9",
        # audio
        "-c:a", "aac", "-b:a", "192k", "-ac", "2",
        str(out)
    ]
    print(">>", " ".join(cmd))
    subprocess.run(cmd, check=True)
print("Conversion done:", len(bad_files), "files fixed.")

```

```
>> ffmpeg -y -i assets\Cosmos_War_of_the_Planets.mp4 -c:v libx264 -preset slow -b:v 3M -vf scale=640:360,fps=25,setdar=16/9 -c:a aac -b:a 192k -ac 2 assets\Cosmos_War_of_the_Planets_formatOK.mp4
>> ffmpeg -y -i assets\Last_man_on_earth_1964.mov -c:v libx264 -preset slow -b:v 3M -vf scale=640:360,fps=25,setdar=16/9 -c:a aac -b:a 192k -ac 2 assets\Last_man_on_earth_1964_formatOK.mp4
>> ffmpeg -y -i assets\The_Gun_and_the_Pulpit.avi -c:v libx264 -preset slow -b:v 3M -vf scale=640:360,fps=25,setdar=16/9 -c:a aac -b:a 192k -ac 2 assets\The_Gun_and_the_Pulpit_formatOK.mp4
>> ffmpeg -y -i assets\Voyage_to_the_Planet_of_Prehistoric_Women.mp4 -c:v libx264 -preset slow -b:v 3M -vf scale=640:360,fps=25,setdar=16/9 -c:a aac -b:a 192k -ac 2 assets\Voyage_to_the_Planet_of_Prehistoric_Women_formatOK.mp4
Conversion done: 4 files fixed.
```

Quick media-format glossary **

Container vs. codec – MP4, MOV, AVI are “boxes”; H.264 and AAC are the payload. A file can be MP4 outside but still hold the wrong codecs.

Frame-rate (fps) – The festival wants **25 fps, the European broadcast standard. A variable or 29.97 fps clip stutters on PAL equipment, so we force `fps=25` .

Resolution & DAR – Scaling to **640×360** keeps the pixels square; `setdar=16/9` tells players the intended display shape.

Bit-rate – `-b:v 3M` is plenty for SD H.264; audio at **192 kb/s stereo** is transparent quality yet small.

`libx264 preset slow` – balances encoding time and quality; “slow” maximises PSNR at this resolution.

`ffprobe` exposes every field in JSON, so the QC script can reject any clip that drifts from these numbers and then `ffmpeg` fixes them in one pass.